

# IRRS: Information Retrieval and Recommender Systems

FIB, Master in Data Science

Slides by Marta Arias, José Luis Balcázar,  
Ramon Ferrer-i-Cancho, Ricard Gavaldá  
Department of Computer Science, UPC

Fall 2022

<http://www.cs.upc.edu/~ir-miri>

## Query answering

A bad algorithm:

```
input query  $q$ ;  
for every document  $d$  in database  
    check if  $d$  matches  $q$ ;  
    if so, add its docid to list  $L$ ;  
output list  $L$  (perhaps sorted in some way);
```

Query processing time should be largely independent of database size.

Probably proportional to answer size.

## 3. Implementation

### Central Data Structure

From terms to documents


A vocabulary or lexicon or dictionary, usually kept in main memory, maintains all the indexed terms (*set*, *map*...); and, besides...

#### The Inverted File


The crucial data structure for indexing.

- ▶ A data structure to support the operation:
  - ▶ “given term  $t$ , get all the documents that contain it”.
- ▶ The inverted file must support this operation (and variants) very efficiently.
- ▶ Built at preprocessing time, not at query time: can afford to spend some time in its construction.


The inverted file: Variant 1



The inverted file: Variant 3



The inverted file: Variant 2



Postings

The inverted file is made of incidence/posting lists

We assign a *document identifier*, docid to each document.  
The dictionary may fit in RAM for medium-size applications.

For each indexed term

a posting list: list of docid's (plus maybe other info) where the term appears.

- ▶ Wonderful if it fits in memory, but this is unlikely.
- ▶ Additionally: posting lists are
  - ▶ almost always sorted by *docid*
  - ▶ often compressed: minimize info to bring from disk!

# Implementation of the Boolean Model, I

Simplest: Traverse posting lists

Conjunctive query:  $a$  AND  $b$


- ▶ intersect the posting lists of  $a$  and  $b$ ;
- ▶ if sorted: can do a merge-like intersection;
- ▶ time: order of the sum of the lengths of posting lists.

```
intersect(input lists  $L1$ ,  $L2$ , output list  $L$ ):  
  while ( not  $L1.end()$  and not  $L2.end()$  )  
    if ( $L1.current() < L2.current()$ )  $L1.advance()$ ;  
    else if ( $L1.current() > L2.current()$ )  $L2.advance()$ ;  
    else {  $L.append(L1.current());$   
           $L1.advance(); L2.advance();$  } }
```

9/31

# Implementation of the Boolean Model, III

Sublinear time intersection: Skip pointers



- ▶ We've merged 1... 19 and 3... 26.
- ▶ We are looking at 36 and 85.
- ▶ Since  $pointer(36) = 62 < 85$ , we can jump to 84 in  $L1$ .

11/31

# Implementation of the Boolean Model, II

Simplest

- ▶ Similar merge-like union for OR.
  - ▶ Time: again order of the sum of lengths of posting lists.
- ▶ Alternative: traverse one list and look up every docid in the other via binary search.
  - ▶ Time: length of shortest list times log of length of longest.


Example:

- ▶  $|L1| = 1000$ ,  $|L2| = 1000$ :
  - ▶ sequential scan: 2000 comparisons,
  - ▶ binary search:  $1000 * 10 = 10,000$  comparisons.
- ▶  $|L1| = 100$ ,  $|L2| = 10,000$ :
  - ▶ sequential scan: 10,000 comparisons,
  - ▶ binary search:  $100 * \log(10,000) = 1400$  comparisons.

10/31

# Implementation of the Boolean Model, IV

Sublinear time intersection: Skip pointers



- ▶ Forward pointer from some elements.
- ▶ Either jump to next segment, or search within next segment (once).
- ▶ Optimal: in RAM,  $\sqrt{|L|}$  pointers of length  $\sqrt{|L|}$ .
- ▶ Difficult to do well, particularly if the lists are on disk.

12/31

Queries can be evaluated according to different plans

E.g.  $a$  AND  $b$  AND  $c$  as

- ▶  $(a$  AND  $b)$  AND  $c$
- ▶  $(b$  AND  $c)$  AND  $a$
- ▶  $(a$  AND  $c)$  AND  $b$

E.g.  $(a$  AND  $b)$  OR  $(a$  AND  $c)$  also as

- ▶  $a$  AND  $(b$  OR  $c)$

The cost of an execution plan depends on the sizes of the lists and the sizes of intermediate lists.

13/31

## Implementation of the Vectorial Model, I

Problem statement

Fixed similarity measure  $sim(d, q)$ :

Retrieve

documents  $d_i$  which have a similarity to the query  $q$

- ▶ either
  - ▶ above a threshold  $sim_{min}$ , or
  - ▶ the top  $r$  according to that similarity, or
  - ▶ all documents,
- ▶ sorted by decreasing similarity to the query  $q$ .

Must react very fast (thus, careful to the interplay with disk!), and with a reasonable memory expense.

15/31

Query:  $(a$  AND  $b)$  OR  $(a$  AND  $c$  AND  $d)$ .

Assume:  $|La| = 3000, |Lb| = 1000, |Lc| = 2500, |Ld| = 300$ .

- ▶ Three intersections plus one union, in the order given: up to cost 13600.
- ▶ Instead,  $((d$  AND  $c)$  AND  $a)$ : reduces to up to cost 11400.
- ▶ Rewrite to  $a$  AND  $(b$  OR  $(c$  AND  $d))$ : reduces to up to cost 8400.

14/31

## Implementation of the Vectorial Model, II

Obvious nonsolution

Traverse all the documents, look at their terms in order to compute similarity, filter according to  $sim_{min}$ , and sort them...

... will not work.

16/31

# Implementation of the Vectorial Model, III

Observations

Most documents include a small proportion of the available terms.  
Queries usually include a humanly small number of terms.  
Only a very small proportion of the documents will be relevant.  
A priori bound  $r$  on the size of the answer known.  
Inverted file available!

17/31

## Index compression, I

Why?

A large part of the query-answering time is spent  
bringing posting lists from disks to RAM.

Need to minimize amount of bits to transfer.

Index compression schemes use:

- ▶ Docid's sorted in increasing order.
- ▶ Frequencies usually very small numbers.
- ▶ Can do better than e.g. 32 bits for each.

19/31

# Implementation of the Vectorial Model, IV

Idea

Invert the loops:

- ▶ Outer loop on the terms  $t$  that appear in the query.
- ▶ Inner loop on documents that contain term  $t$ .
  - ▶ the reason for inverted index!
- ▶ Accumulate similarity for visited documents.
- ▶ Upon termination, normalize and sort.

Many additional subtleties can be incorporated.

18/31

## Index compression, II

Why?

A large part of the query-answering time is spent bringing  
posting lists from disks to RAM.

- ▶ Need to minimize amount of bits to transfer.

Easiest is to use “`int` type” to store docid's and frequencies

- ▶ 8 bytes, 64 bits per pair
- ▶ ... but want/can/need to do much better!

Index compression schemes use:

- ▶ Docid's sorted in increasing order.
- ▶ Frequencies usually very small numbers.

20/31

Posting list is:

$$term \rightarrow [(id_1, f_1), (id_2, f_2), \dots, (id_k, f_k)]$$

Can we compress frequencies  $f_i$ ?:

Yes! Will use *unary self-delimiting* codes because frequencies typically very small

Can we compress docid's  $id_i$ ?:

Yes! Will use *Gap compression* and *Elias Gamma* codes because docid's are sorted

21/31

## Index compression, V

Compressing frequencies: unary encoding

Unary encoding of  $x$  is  $\overbrace{111 \dots 1}^{x \text{ times}}$

- E.g.  $unary(15) = 1111111111111111$
- $|unary(x)| = x$ 
  - typical binary encoding:  $|binary(x)| = \log_2(x)$
- variable length encoding

But..

want to encode *lists* of frequencies, where do we cut?

23/31

The distribution of frequencies is very biased towards small numbers, i.e., most  $f_i$  are very small

- Exercise: can you quantify this using Zipf's law?
- E.g. in files for lab session 1: 68 % is 1, 13 % is 2, 6 % is 3, <13 % is >3, <3 % is >10, 0.6 % is >20.

Unary code

Want encoding scheme that uses few bits for small frequencies

22/31

## Index compression, VI

Compressing frequencies: self-delimiting unary encoding

- Make 0 act as a separator
- Replace last 1 in each number with a 0
- Example: [3, 2, 1, 4, 1, 5] encoded as 110 10 0 1110 0 11110
- This is a *self-delimiting* code: no prefix of a code is a code
- Self-delimiting *implies* unique decoding

24/31

# Index compression, VII

Compressing frequencies: self-delimiting unary encoding

Recall example from lab session 1: 68 % is 1, 13 % is 2, 6 % is 3, <13 % is >3, <3 % is >10, 0.6 % is >20, the expected length would be (approx)

$$1 * 0.68 + 2 * 0.13 + 3 * 0.06 + 6^1 * 0.13 = 1.91$$

Unary code works very well

- ▶ 1 bit when  $f_i = 1$
- ▶ 1.3 to 2.5 bits per  $f_i$  on real corpuses
- ▶ 1 bit per term occurrence in document
  - ▶ Easy to estimate memory used!

---

<sup>1</sup>I put it something greater than 3 as an approximation

25 / 31

# Index compression, IX

Compressing docid's

- ▶ Fewer bits if gaps are small
- ▶ E.g.:  $N = 10^6$ ,  $|L| = 10^4$ , then average gap is 100
  - ▶ So, could use 8 bits instead of 20 (or 32)
- ▶ .. but .. this is only on average! *Large gaps do exist*
  - ▶ Will need a *variable length, self-delimiting* encoding scheme
- ▶ Gaps are not biased towards 1, so unary not a good idea
  - ▶ Will use need a *variable length, self-delimiting, binary* encoding scheme

27 / 31

# Index compression, VIII

Compressing docid's

Gap compression

Instead of compressing  $[(id_1, f_1), (id_2, f_2), \dots, (id_k, f_k)]$   
Compress  $[(id_1, f_1), (id_2 - id_1, f_2), \dots, (id_k - id_{k-1}, f_k)]$

Example:

$(1000, 1), (1021, 2), (1037, 1), (1056, 4), (1080, 1), (1095, 3)$   
compressed to:

$(1000, 1), (21, 2), (16, 1), (19, 4), (24, 1), (15, 3)$

26 / 31

# Index compression, X

Compressing docid's: Elias-Gamma code (self-delimiting binary code)

IDEA:

First say how long  $x$  is in binary, then send  $x$

Pseudo-code for Elias-Gamma encoding:

- ▶ let  $w = \text{binary}(x)$
- ▶ let  $y = |w|$
- ▶ prepend  $y - 1$  zeros to  $w$ , and return

Examples:

$EG(1) = 1, EG(2) = 010, EG(3) = 011, EG(4) = 00100, EG(20) = 000010100$

28 / 31

## Index compression, XI

Compressing docid's: Elias-Gamma code (self-delimiting binary code)

- ▶ Elias-Gamma code is self-delimiting
  - ▶ Exercise: think how to decode uniquely
- ▶ Length of a code for  $x$  is about  $2 \log_2(x)$ 
  - ▶ Exercise: why?

29 / 31

## Index compression, XII

Compressing docid's: easier alternative, *variable byte codes*

Easier alternative: byte-wise (8 bits) or nibble-wise (4 bits) encoding that make use of first bit to say whether it is the last byte or not (*continuation bit*).

- ▶ Encoding is also variable length, but much simpler
- ▶ Waste is not that much
- ▶ Better use of CPU by reading bytes instead of single bits
- ▶ First bit of byte is continuation bit, other 7 bits used to encode in binary
  - ▶ if 0, then last byte
  - ▶ if 1, number continues

Example:

10101001 11100111 01100111 is code for  
0101001 1100111 1100111 (continuation bits in red)

30 / 31

## Index compression, XIII

Bottom line

- ▶ Ratios of 20 % to 25 % routinely achieved
- ▶ Translates to similar speed-up at query time

31 / 31