

CAIM: Cerca i Anàlisi d'Informació Massiva

FIB, Grau en Enginyeria Informàtica

Slides by Marta Arias, José Luis Balcázar,
Ramon Ferrer-i-Cancho, Ricard Gavaldà
Department of Computer Science, UPC

Fall 2018

<http://www.cs.upc.edu/~caim>

Searching the Web, I

When documents are interconnected

The World Wide Web is huge


- ▶ 100,000 indexed pages in 1994
- ▶ 10,000,000,000's indexed pages in 2013
- ▶ Most queries will return *millions* of pages with high similarity.
- ▶ Content (text) alone cannot discriminate.
- ▶ Use the structure of the Web - a graph.
- ▶ Gives indications of the prestige - usefulness of each page.

5. Web Search. Architecture of simple IR systems

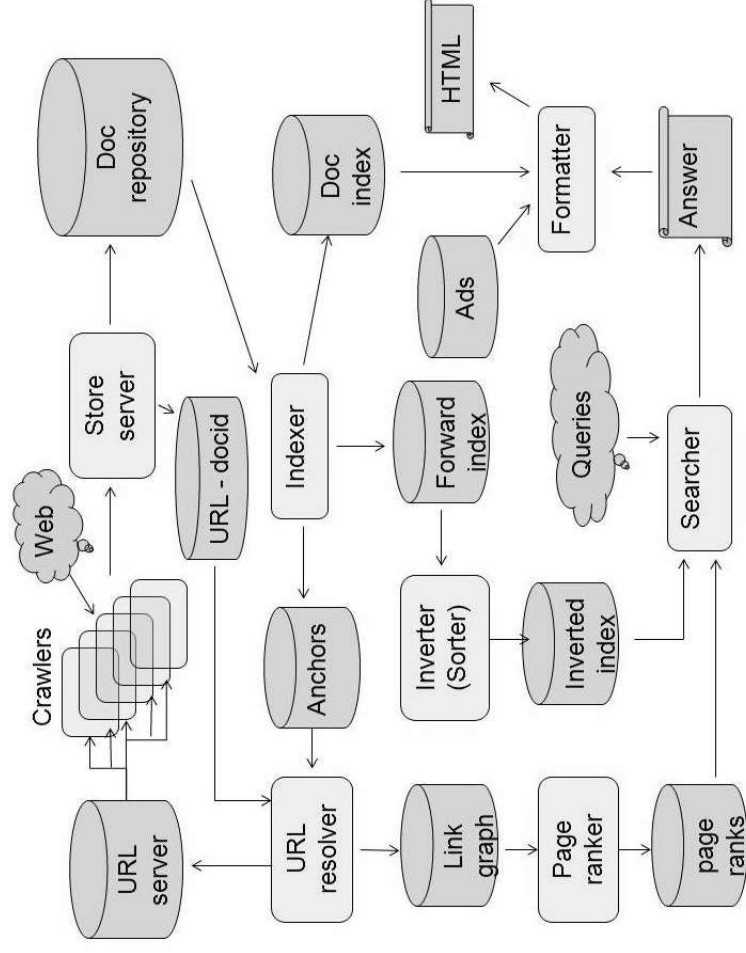
How Google worked in 1998

S. Brin, L. Page: "The Anatomy of a Large-Scale Hypertextual Web Search Engine", 1998

Notation:



Some components



The inverter (sorter), I

Transforms forward index to inverted index

First idea:

```

for every entry document d
  for every term t in d
    add docid(d) at end of list for t;

```

Lousy locality, many disk seeks, too slow

7/44

Better idea for indexing:

```

create in disk an empty inverted file, ID;
create in RAM an empty index IR;
for every document d
  for every term t in d
    add docid(d) at end of list for t in IR;
  if RAM full
    for each t, merge the list for t in IR
      into the list for t in ID;

```

Merging previously sorted lists is sequential access
Much better locality. Much fewer disk seeks.


The inverter (sorter), II

6/44

8/44

The inverter (sorter), III

The above can be done concurrently on different sets of documents:



9 / 44

10 / 44

Searching the Web, I


When documents are interconnected

The internet is huge

- ▶ 100,000 indexed pages in 1994
- ▶ 10,000,000,000 indexed pages at end of 2011

To find content, it is necessary to search for it

- ▶ We know how to deal with the content of the webpages
- ▶ But.. what can we do with the structure of the internet?



11 / 44

The inverter (sorter), IV

- ▶ Indexer ships barrels, fragments of forward index
- ▶ Barrel size = what fits in main memory
- ▶ Separately, concurrently inverted in main memory
- ▶ Inverted barrels merged to inverted index
- ▶ 1 day instead of estimated months

Searching the Web, II

Meaning of a hyperlink

When page A links to page B , this means

- ▶ A 's author thinks that B 's content is interesting or important
- ▶ So a link from A to B , adds to B 's reputation


But not all links are equal..

- ▶ If A is very important, then $A \rightarrow B$ "counts more"
- ▶ If A is not important, then $A \rightarrow B$ "counts less"

In today's lecture we'll see two algorithms based on this idea

- ▶ *Pagerank* (Brin and Page, oct. 98)
- ▶ *HITS* (Kleinberg, apr. 98)

12 / 44



Pagerank, I

The idea that made Google great

Pagerank, II

Definitions

The web is a graph $G = (V, E)$

- ▶ $V = \{1, \dots, n\}$ are the nodes (that is, the pages)
- ▶ $(i, j) \in E$ if page i points to page j
- ▶ we associate to each page i , a real value p_i (i 's pagerank)
- ▶ we impose that $\sum_{i=1}^n p_i = 1$

Intuition:

A page is important if it is pointed to by other important pages

- ▶ Circular definition ...
- ▶ not a problem!

How are the p_i 's related

- ▶ p_i depends on the values p_j of pages j pointing to i

$$p_i = \sum_{j \rightarrow i} \frac{p_j}{out(j)}$$


- ▶ where $out(j)$ is j 's outdegree

13 / 44

14 / 44

Pagerank, III

Example



A set of $n + 1$ linear equations:

$$\begin{aligned} p_1 &= \frac{p_1}{3} + \frac{p_2}{2} \\ p_2 &= \frac{p_3}{2} + p_4 \\ p_3 &= \frac{p_1}{3} \\ p_4 &= \frac{p_1}{3} + \frac{p_2}{2} + \frac{p_3}{2} \\ 1 &= p_1 + p_2 + p_3 + p_4 \end{aligned}$$

Whose solutions is:

$$p_1 = 6/23, p_2 = 8/23, p_3 = 2/23, p_4 = 7/23$$

15 / 44

Pagerank, IV

Formally

Equations

- ▶ $p_i = \sum_{j: (j,i) \in E} \frac{p_j}{out(j)}$ for each $i \in V$
- ▶ $\sum_{i=1}^n p_i = 1$

where $out(i) = |\{j : (i, j) \in E\}|$ is the outdegree of node i

If $|V| = n$


- ▶ $n + 1$ equations
- ▶ n unknowns

Could be solved, for example, using Gaussian elimination in time $O(n^3)$

16 / 44

Pagerank, V

Example, revisited



A set of linear equations:

$$\begin{pmatrix} p_1 \\ p_2 \\ p_3 \\ p_4 \end{pmatrix} = \begin{pmatrix} \frac{1}{3} & \frac{1}{2} & 0 & 0 \\ 0 & 0 & \frac{1}{2} & 1 \\ \frac{1}{3} & 0 & 0 & 0 \\ \frac{1}{3} & \frac{1}{2} & \frac{1}{2} & 0 \end{pmatrix} \cdot \begin{pmatrix} p_1 \\ p_2 \\ p_3 \\ p_4 \end{pmatrix}$$

namely: $\vec{p} = M^T \vec{p}$ and additionally $\sum_i p_i = 1$


Whose solutions is:

\vec{p} is the eigenvector of matrix M^T associated to eigenvalue 1

17 / 44

Pagerank, VII

Example, revisited



$$M = \begin{pmatrix} 1/3 & 0 & 1/3 & 1/3 \\ 1/2 & 0 & 0 & 1/2 \\ 0 & 1/2 & 0 & 1/2 \\ 0 & 1 & 0 & 0 \end{pmatrix}$$

(rows add up to 1)


$$M^T = \begin{pmatrix} 1/3 & 1/2 & 0 & 0 \\ 0 & 0 & 1/2 & 1 \\ 1/3 & 0 & 0 & 0 \\ 1/3 & 1/2 & 1/2 & 0 \end{pmatrix}$$

(columns add up to 1)

19 / 44

Pagerank, VI

Example, revisited



What does M^T look like?

$$M^T = \begin{pmatrix} \frac{1}{3} & \frac{1}{2} & 0 & 0 \\ 0 & 0 & \frac{1}{2} & 1 \\ \frac{1}{3} & 0 & 0 & 0 \\ \frac{1}{3} & \frac{1}{2} & \frac{1}{2} & 0 \end{pmatrix}$$

M^T is the *transpose* of the row-normalized adjacency matrix of the graph !

18 / 44

Pagerank, VIII

Example, revisited

$$A = \begin{pmatrix} 1 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 \end{pmatrix} \quad M = \begin{pmatrix} \frac{1}{3} & \frac{1}{2} & 0 & \frac{1}{3} \\ \frac{1}{2} & 0 & \frac{1}{2} & 1 \\ 0 & \frac{1}{2} & 0 & \frac{1}{2} \\ 0 & 1 & 0 & 0 \end{pmatrix} \quad M^T = \begin{pmatrix} \frac{1}{3} & 0 & \frac{1}{3} & \frac{1}{3} \\ 0 & 0 & \frac{1}{2} & 1 \\ \frac{1}{3} & \frac{1}{2} & 1 & \frac{1}{2} \\ 0 & 0 & 0 & 0 \end{pmatrix}$$

Question:

Why do we need to row-normalize and transpose A?

Answer:

- **Row normalization:** because $p_i = \sum_{j:(j,i) \in E} \frac{p_j}{out(j)}$
- **Transpose:** because $p_i = \sum_{j:(j,i) \in E} \frac{p_j}{out(j)}$, that is, p_i depends on i 's *incoming edges*

20 / 44



Pagerank, IX

It is just about solving a system of linear equations!

.. but

- ▶ How do we know a solution exists?
- ▶ How do we know it has a single solution?
- ▶ How can we compute it efficiently?

For example, the graph on the left has no solution.. (check it!) but the one on the right does




21 / 44

Pagerank, XI

Equivalently: the random surfer view

Now assume M is the transition probability matrix between states in G



$$M = \begin{pmatrix} 1/3 & 0 & 1/3 & 1/3 \\ 1/2 & 0 & 0 & 1/2 \\ 0 & 1/2 & 0 & 1/2 \\ 0 & 1 & 0 & 0 \end{pmatrix}$$

Let $\vec{p}(t)$ be the probability over states at time t

- ▶ E.g., $p_j(0)$ is the probability of being at state j at time 0

Random surfer jumps from page i to page j with probability m_{ij}

- ▶ E.g., probability of transitioning from state 2 to state 4 is $m_{24} = 1/2$

23 / 44

Pagerank, X

How do we know a solution exists?

Luckily, we have some results from *linear algebra*

Definition

A matrix M is stochastic, if

- ▶ All entries are in the range $[0, 1]$
- ▶ Each row adds up to 1 (i.e., M is row normalized)

Theorem (Perron-Frobenius)

If M is stochastic, then it has at least one stationary vector, i.e., one non-zero vector p such that

$$M^T p = p.$$

22 / 44

Pagerank, XII

The random surfer view

- ▶ Surfer starts at random page according to probability distribution $\vec{p}(0)$
- ▶ At time $t > 0$, random surfer follows one of current page's links uniformly at random

$$\vec{p}(t) := M^T \vec{p}(t-1)$$


- ▶ In the limit $t \rightarrow \infty$:

- ▶ $\vec{p}(t) = \vec{p}(t+1) = \vec{p}(t+2) = \dots = \vec{p}$
- ▶ so $\vec{p}(t) = M^T \vec{p}(t-1)$
- ▶ $\vec{p}(t)$ converges to a solution p s.t. $p = M^T p$ (the pagerank solution)!

24 / 44

Pagerank, XIII

Random surfer example




$$M^T = \begin{pmatrix} \frac{1}{3} & \frac{1}{2} & 0 & \frac{1}{3} \\ 0 & \frac{1}{2} & 0 & \frac{1}{2} \\ \frac{1}{3} & 0 & 0 & \frac{1}{2} \\ \frac{1}{3} & \frac{1}{2} & 1 & 0 \end{pmatrix}$$

- ▶ $\vec{p}(0)^T = (1, 0, 0, 0)$
- ▶ $\vec{p}(1)^T = (1/3, 0, 1/3, 1/3)$
- ▶ $\vec{p}(2)^T = (0.11, 0.50, 0.11, 0.28)$
- ▶ ..
- ▶ $\vec{p}(10)^T = (0.26, 0.35, 0.09, 0.30)$
- ▶ $\vec{p}(11)^T = (0.26, 0.35, 0.09, 0.30)$

25 / 44

Pagerank, XV

Convergence of the Power method: aperiodicity required



Try out the power method with $\vec{p}(0)$:

$$\begin{pmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{pmatrix}, \text{ or } \begin{pmatrix} 1 \\ 0 \\ 0 \\ 0 \end{pmatrix}, \text{ or } \begin{pmatrix} 1/2 \\ 0 \\ 1/2 \\ 0 \end{pmatrix}$$

Not being able to break the cycle looks problematic!

- ▶ .. so will require graphs to be aperiodic
 - ▶ no integer $k > 1$ dividing the length of every cycle

27 / 44

Pagerank, XIV

An algorithm to solve the eigenvector problem (find p s.t. $p = M^T p$)

The Power Method

- ▶ Chose initial vector $\vec{p}(0)$ randomly
- ▶ Repeat $\vec{p}(t) \leftarrow M^T \vec{p}(t-1)$
- ▶ Until convergence (i.e. $\vec{p}(t) \approx \vec{p}(t-1)$)

We are hoping that


- ▶ The method converges
- ▶ The method converges fast
- ▶ The method converges fast to the pagerank solution
- ▶ The method converges fast to the pagerank solution regardless of the initial vector

26 / 44

Pagerank, XVI

Convergence of the Power method: strong connectedness required

What happens with the pagerank in this graph?



The sink hoards all the pagerank!

- ▶ need a way to leave sinks
- ▶ .. so we will force graphs to be strongly connected

28 / 44

Pagerank, XVII

A useful theorem from Markov chain theory

Theorem

If a matrix M is strongly connected and aperiodic, then:

- ▶ $M^T \vec{p} = \vec{p}$ has exactly one non-zero solution such that $\sum_i p_i = 1$
- ▶ 1 is the largest eigenvalue of M^T
- ▶ the Power method converges to the \vec{p} satisfying $M^T \vec{p} = \vec{p}$, from any initial non-zero $\vec{p}(0)$
- ▶ Furthermore, we have exponential fast convergence

To guarantee a solution, we will make sure that the matrices that we work with are strongly connected and aperiodic

29 / 44

Pagerank, XIX

Teleportation in the random surfer view

The meaning of λ

- ▶ With probability λ , the random surfer follows a link in current page
- ▶ With probability $1 - \lambda$, the random surfer jumps to a random page in the graph (teleportation)

31 / 44

Pagerank, XVIII

Guaranteeing aperiodicity and strong connectedness

Definition (The Google Matrix)

Given a damping factor λ such that: $0 < \lambda < 1$:

$$G = \lambda M + (1 - \lambda) \frac{1}{n} J$$

where J is a $n \times n$ matrix containing 1 in each entry

Observe that:


- ▶ G is stochastic
 - ▶ .. because G is a weighted average of M and $\frac{1}{n} J$, which are also stochastic
- ▶ for each integer $k > 0$, there is a non-zero probability path of length k from every state to any other state of G
 - ▶ .. implying that G is strongly connected and aperiodic
- ▶ and so the Power method will converge on G , and fast!

30 / 44

Pagerank, XX

Exercise, I

Compute the pagerank value of each node of the following graph assuming a damping factor $\lambda = 2/3$:



Hint: solve the following system, using $p_2 = p_3 = p_4$

$$\begin{pmatrix} p_1 \\ p_2 \\ p_3 \\ p_4 \end{pmatrix} = \begin{bmatrix} 0 & 1 & 1 & 1 \\ \frac{2}{3} & \frac{1}{3} & 0 & 0 \\ \frac{2}{3} & \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \\ \frac{2}{3} & \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \end{bmatrix} \begin{pmatrix} p_1 \\ p_2 \\ p_3 \\ p_4 \end{pmatrix} + \frac{1}{3} \cdot \frac{1}{4} \begin{pmatrix} 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \end{pmatrix} \begin{pmatrix} p_1 \\ p_2 \\ p_3 \\ p_4 \end{pmatrix}$$

32 / 44

Compute the pagerank vector \vec{p} of graph with row-normalized matrix M for damping factor λ in closed matrix form.

Answer:

$$\vec{p} = (I - \lambda M^T)^{-1} \begin{pmatrix} \frac{1-\lambda}{n} \\ \vdots \\ \frac{1-\lambda}{n} \end{pmatrix}$$

33 / 44

Topic-sensitive Pagerank, II

Assume there is a small set of K topics (sports, science, politics, ...)

- ▶ Each topic $k \in \{1, \dots, K\}$ is defined by a subset of the web pages T_k
- ▶ For each k , compute pagerank of node i for topic k :

$p_{i,k}$ = “pagerank of node i with teleportation reduced to T_k ”

- ▶ Finally compute ranking score of a page i given query q

$$score(i, q) = \sum_{k=1}^K sim(T_k, q) \cdot p_{i,k}$$

35 / 44

Topic-sensitive Pagerank, I

Observe that pageranks are independent of user's query

- ▶ Advantages
 - ▶ Computed off-line
 - ▶ Collective reputation
- ▶ Disadvantages
 - ▶ Insensitive to particular user's needs

34 / 44

HITS, I

Hypertext Induced Text Search

Interest of a web page due to two different reasons

- ▶ page content is interesting (*authority*), or
- ▶ page points to interesting pages (*hub*)

HITS main rationale

- ▶ hubs are important if they point to important authorities
- ▶ authorities are important if pointed to by important hubs
- ▶ .. but .. circular definition again not a problem!

36 / 44

HITS, II

Definition of authority and hub value (a_i and h_i)

Associate to each page i an authority value a_i and a hub value h_i

- ▶ vector of all authority values is \vec{a}
- ▶ vector of all hub values is \vec{h}

Keep these vectors normalized (notice L2 norm!)

- ▶ $\|\vec{a}\| = \sum_i a_i^2 = 1$, and $\|\vec{h}\| = \sum_i h_i^2 = 1$

For appropriate scaling constants c and d

- ▶ $a_i = c \cdot \sum_{j \rightarrow i} h_j$, and $h_i = d \cdot \sum_{i \rightarrow j} a_j$

Notice not a linear system anymore!


- ▶ ... but still ok with a variant of the power method

37 / 44

HITS, IV

Example

Our old graph



Adjacency matrix

$$A = \begin{pmatrix} 1 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 \end{pmatrix}$$

$h_2 = d \cdot (a_1 + a_4)$ // here we use A 's second row


$$h_2 \propto (1, 0, 0, 1) \cdot \begin{pmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \end{pmatrix} = (1, 0, 0, 1) \cdot \vec{a}$$

39 / 44

HITS, III

Example

Our old graph



Adjacency matrix

$$A = \begin{pmatrix} 1 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 \end{pmatrix}$$

$a_1 = c \cdot (h_1 + h_2)$ // here we use A 's first column

$$a_1 \propto (1, 1, 0, 0) \cdot \begin{pmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \end{pmatrix} = (1, 1, 0, 0) \cdot \vec{h}$$

38 / 44

HITS, V

Update rule for \vec{a} and \vec{h}

Written in compact matrix form

- ▶ To update authority values
 - ▶ $\vec{a} := A^T \cdot \vec{h}$
 - ▶ normalize afterwards $\vec{a} := \frac{\vec{a}}{\|\vec{a}\|}$ so that $\|\vec{a}\| = 1$
- ▶ To update hub values
 - ▶ $\vec{h} := A \cdot \vec{a}$
 - ▶ normalize afterwards $\vec{h} := \frac{\vec{h}}{\|\vec{h}\|}$ so that $\|\vec{h}\| = 1$

40 / 44

HITS, VI

The power method for finding \vec{a} and \vec{h}

Given adjacency matrix A

- ▶ Initialize $\vec{a} = \vec{h} = (1, 1, \dots, 1)^T$
- ▶ Normalize \vec{a} and \vec{h} so that $\|a\| = \|h\| = 1$
- ▶ Repeat until convergence
 - ▶ $\vec{a} := A^T \cdot \vec{h}$
 - ▶ normalize \vec{a} so that $\|a\| = 1$
 - ▶ $\vec{h} := A \cdot \vec{a}$
 - ▶ normalize \vec{h} so that $\|h\| = 1$

41 / 44

HITS, VIII

HITS algorithm illustrated




FIG. 1. Expanding the root set into a base set.

43 / 44

HITS, VII

HITS algorithm

Query answering algorithm HITS

- ▶ Get query q and run content-based searcher on q
- ▶ Let $RootSet$ be the top- k ranked pages
- ▶ Expand pages to $BaseSet$ by adding all pages pointed to and by pages in $RootSet$
- ▶ compute hub and authority values for the subgraph of web induced by $BaseSet$
- ▶ Rank pages in $BaseSet$ according to \vec{a} , \vec{h} , and content

42 / 44

HITS vs. Pagerank

Pros of HITS vs. Pagerank

- ▶ Sensitive to user queries

Cons of HITS vs. Pagerank

- ▶ Compute online, not offline!
- ▶ More vulnerable to webspamming

44 / 44